

Record Label Analysis

01

**Business
Value**

02

Data

03

**Summary
Plots**



04

Predictive Models

05

Comparison of Performance

06

Conclusion

Why Care ?

Gain insight into what kind of artists a record label should sign/put money into by assessing the type of music that performed well in the year 2019. Models will test the relevance certain variables have when predicting popularity and a hit.

Introduction to Data

- **Data:**
 - List of most popular songs from 1950-2019
 - Sentiment Analysis
 - Two Columns Added:
 - Popularity - Spotify API
 - Streams - Manually (Youtube API had vague restrictions)

```
0%matplotlib inline
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()

# Load the data
url = "https://raw.githubusercontent.com/jbrownlee/Datasets/master/spotify.csv"
data = pd.read_csv(url)
print(data.info())

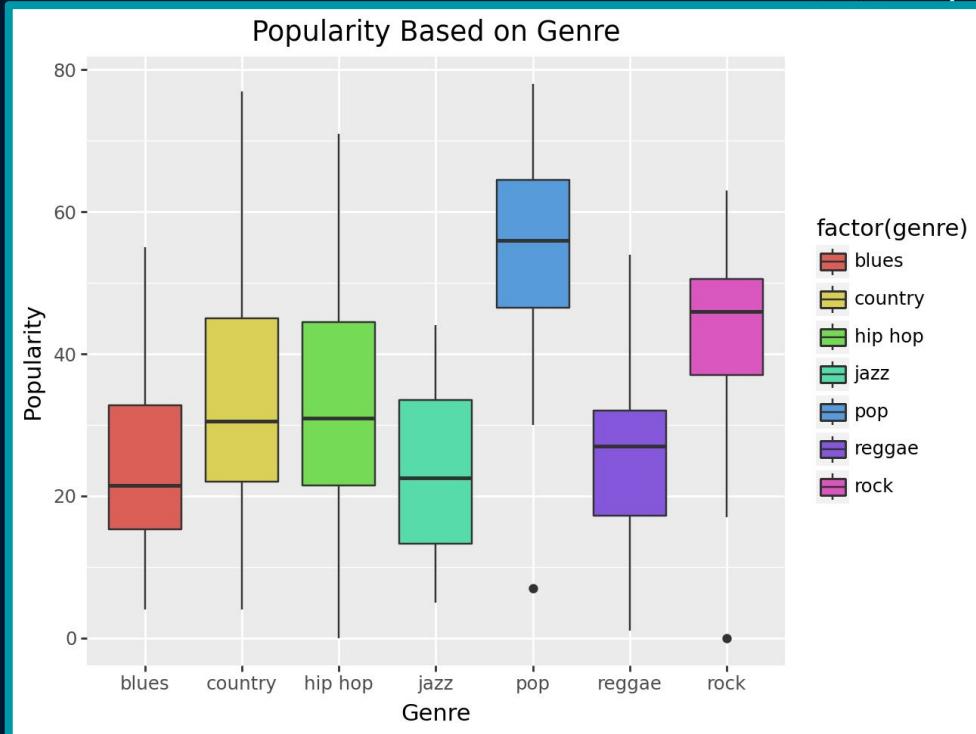
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 531 entries, 0 to 530
Data columns (total 34 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   index            531 non-null    int64  
 1   Unnamed: 0        531 non-null    int64  
 2   artist_name      531 non-null    object  
 3   track_name       531 non-null    object  
 4   release_date     531 non-null    int64  
 5   genre            531 non-null    object  
 6   lyrics           531 non-null    object  
 7   len               531 non-null    int64  
 8   dating            531 non-null    float64 
 9   violence          531 non-null    float64 
 10  world/life        531 non-null    float64 
 11  night/time        531 non-null    float64 
 12  shake the audience 531 non-null    float64 
 13  family/gospel     531 non-null    float64 
 14  romantic          531 non-null    float64 
 15  communication     531 non-null    float64 
 16  obscene           531 non-null    float64 
 17  music              531 non-null    float64 
 18  movement/places   531 non-null    float64 
 19  light/visual perceptions 531 non-null    float64 
 20  happy/spiritual    531 non-null    float64 
 21  like/girls         531 non-null    float64 
 22  sadness            531 non-null    float64 
 23  feelings           531 non-null    float64 
 24  danceability       531 non-null    float64 
 25  loudness           531 non-null    float64 
 26  acoustiveness      531 non-null    float64 
 27  instrumentalness   531 non-null    float64 
 28  valence            531 non-null    float64 
 29  energy              531 non-null    float64 
 30  topic               531 non-null    object  
 31  age                 531 non-null    float64 
 32  popularity          531 non-null    float64 
 33  streams             531 non-null    float64 
dtypes: float64(25), int64(4), object(5)
memory usage: 141.2+ KB
None
```

Summary of Data Cleaning/Feature Engineering

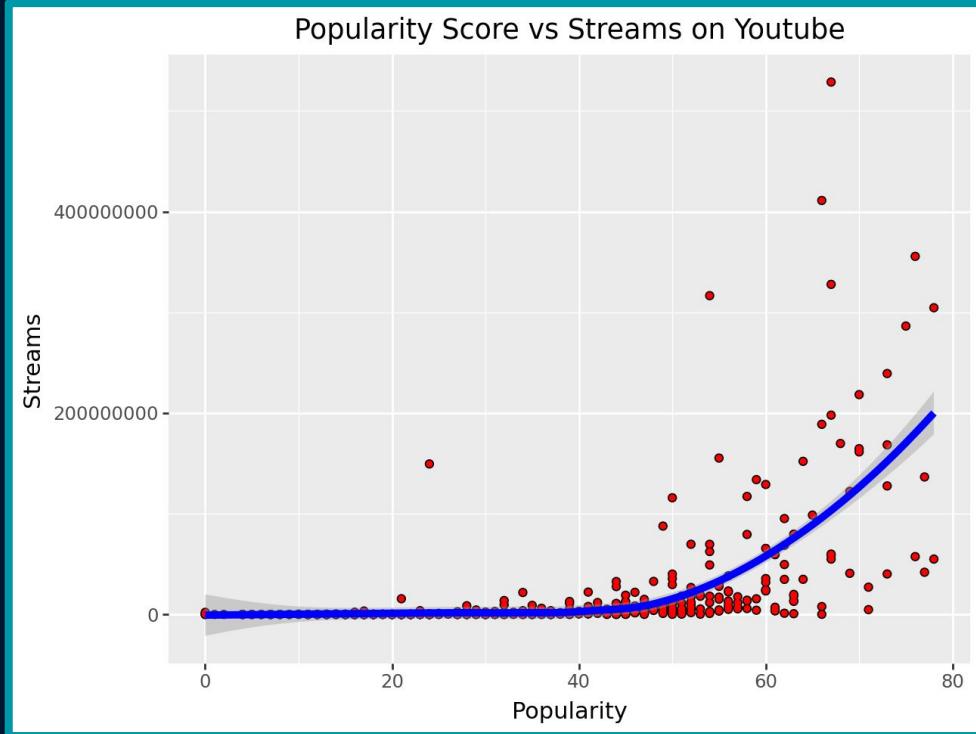
1. Cut all data from songs not in 2019
2. Separate Features from Popularity
(Hit for Logistic)
 - a. $\geq 1mn$
3. Encode Categorical, Scale
Continuous
4. Transform Features
5. Split Data



Summary Plot 1



Summary Plot 2





Models

Beta Regression, Clustering, ElasticNet, PCA

Clustering Model

- **What is Clustering?**
 - Find groups in data
- How to determine best type of model?

```
Make ggplot scatterplots of pairs of your features to give you a little bit of information about the data, and to help you decide which algorithm to use (you don't need to make scatterplots for all possible pairs of features, just make sure each feature appears at least once).

from plotnine import ggplot, aes, geom_point, labs, scale_x_log10, scale_y_log10, theme_minimal

columns = ['genre', 'len', 'dating', 'violence', 'world/life', 'night/time', 'shake the audience',
           'family/gospel', 'romantic', 'communication', 'obscene', 'music', 'movement/places',
           'light/visual perceptions', 'family/spiritual', 'like/girls', 'sadness', 'feelings',
           'danceability', 'loudness', 'acousticness', 'instrumentalness', 'valence', 'energy',
           'topic', 'age', 'popularity', 'streams']

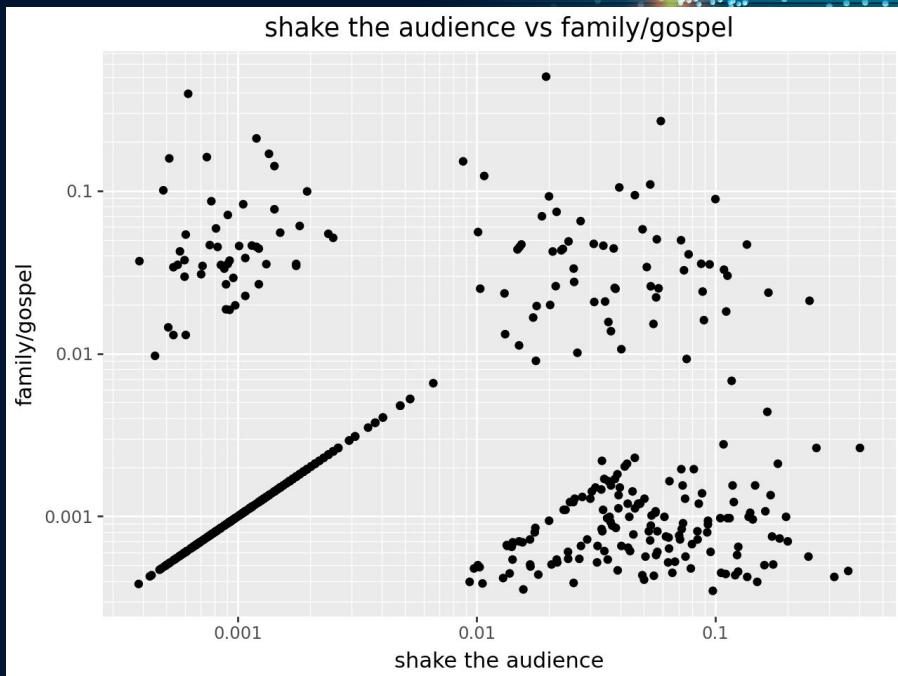
for i in range(0, len(columns) - 1, 2):
    col_x = columns[i]
    col_y = columns[i + 1]

    # Skip if either column contains categorical data or invalid values for log scale CHATGPT
    if data[col_x].dtype == 'object' or data[col_y].dtype == 'object':
        print(f"Skipping plot for {col_x} vs {col_y} due to categorical data.")
        continue

    if (data[col_x] <= 0).any() or (data[col_y] <= 0).any():
        print(f"Skipping plot for {col_x} vs {col_y} due to non-positive values.")
        continue

    plot = (ggplot(data, aes(x=col_x, y=col_y)) +
            geom_point() +
            labs(title=f'{col_x} vs {col_y}', x=col_x, y=col_y) +
            scale_x_log10() + # Log scale for x-axis
            scale_y_log10()) # Log scale for y-axis

    # Print each plot
    display(plot)
```



Clustering Model

- GMM
 - o Multiple clusters in data (6)
 - o Probabilistic
 - o Variances can differ
- Log Likelihood
 - o 20.01 (not great)

```
gmm = GaussianMixture(n_components=5)

# 3. Fit model + predict
labels = gmm.fit_predict(X_processed)

# Add cluster labels to the original DataFrame
data["clusters"] = labels

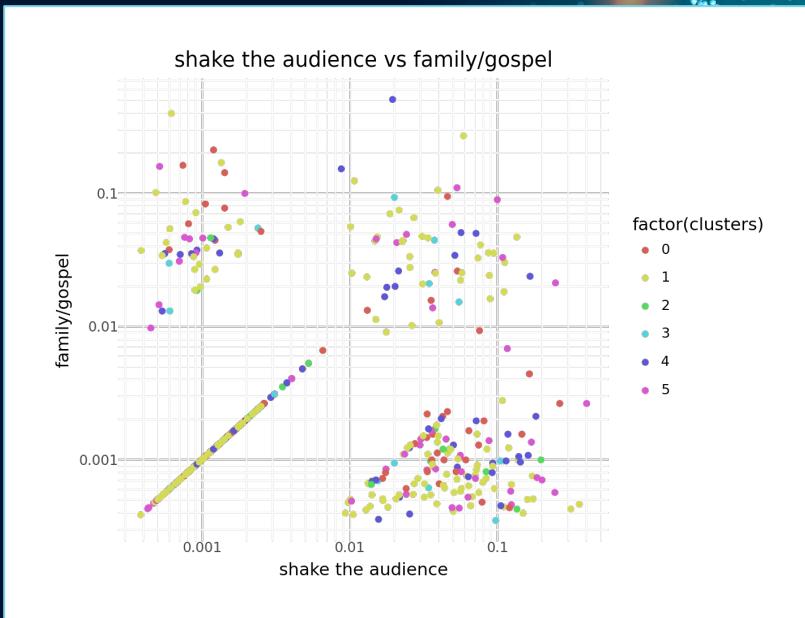
for i in range(0, len(columns) - 1, 2):
    col_x = columns[i]
    col_y = columns[i + 1]

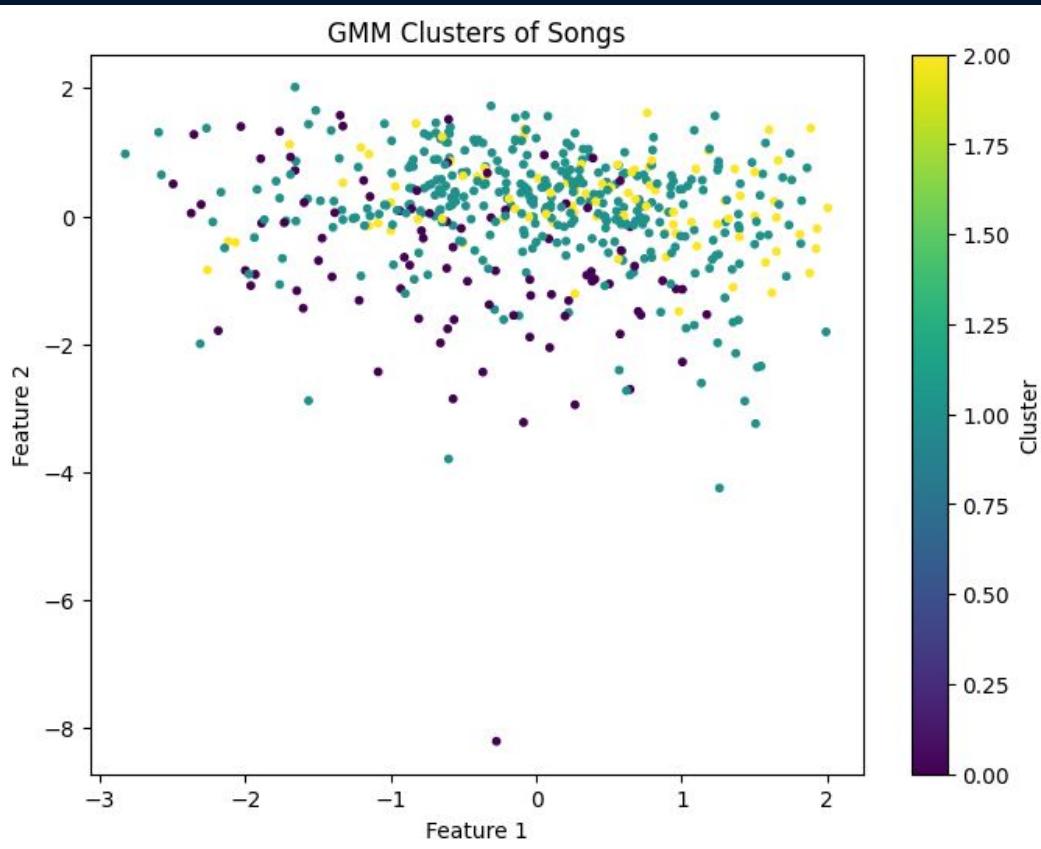
    # Skip if either column contains categorical data or invalid values for log scale
    if data[col_x].dtype == 'object' or data[col_y].dtype == 'object':
        print(f"Skipping plot for {col_x} vs {col_y} due to categorical data.")
        continue

    if (data[col_x] <= 0).any() or (data[col_y] <= 0).any():
        print(f"Skipping plot for {col_x} vs {col_y} due to non-positive values.")
        continue

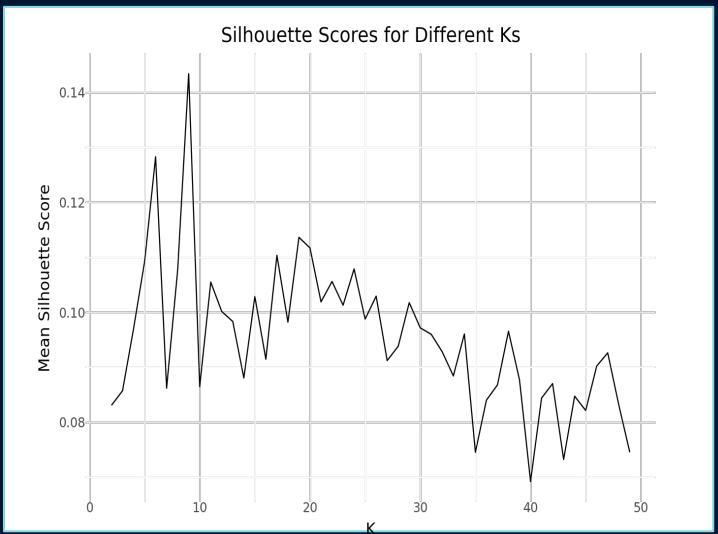
    # Scatter plot
    plot = (ggplot(data, aes(x=col_x, y=col_y, color='factor(clusters)')) +
            geom_point() +
            labs(title=f'{col_x} vs {col_y}', x=col_x, y=col_y) +
            scale_x_log10() + # Log scale for x-axis
            scale_y_log10()) # Log scale for y-axis

    # Print each plot
    plot
```





Avg. cluster means for popularity



```
#predictors  
  
db = DBSCAN(eps = 0.1, min_samples = 100)  
  
# fit  
labels = db.fit_predict(X_processed)  
  
# Add cluster labels to df  
X["clusters"] = labels  
  
silhouette_avg = silhouette_score(X_processed, labels)  
print(f"Silhouette Score: {silhouette_avg}")  
  
Unique labels: {-1}
```

DBSCAN and KMeans Results

Beta Regression

- Why beta regression? (What is beta regression)
- Necessary transformations made to “popularity” column in order to fit the requirements of a beta regression model
- Consistent performance between training and testing sets indicates the model is properly fit and can generalize well on new (unseen) data

```
scaler = MinMaxScaler(feature_range=(1e-6, 1 - 1e-6))
y_scaled = scaler.fit_transform(y.values.reshape(-1, 1)).flatten()
```

```
X_train["streams_log"] = X_train["streams"].apply(lambda x: np.log(x + 1))
X_test["streams_log"] = X_test["streams"].apply(lambda x: np.log(x + 1))
```

```
y_train_pred_original = scaler.inverse_transform(y_train_pred.to_numpy().reshape(-1, 1)).flatten()
y_test_pred_original = scaler.inverse_transform(y_test_pred.to_numpy().reshape(-1, 1)).flatten()

y_train_original = scaler.inverse_transform(y_train.reshape(-1, 1)).flatten()
y_test_original = scaler.inverse_transform(y_test.reshape(-1, 1)).flatten()
```

Generalized Linear Model Regression Results							
Dep. Variable:	y_train	No. Observations:	424				
Model:	GLM	Df Residuals:	422				
Model Family:	Binomial	Df Model:	1				
Link Function:	Logit	Scale:	1.0000				
Method:	IRLS	Log-Likelihood:	-179.94				
Date:	Sat, 30 Nov 2024	Deviance:	33.129				
Time:	23:10:54	Pearson chi2:	32.3				
No. Iterations:	4	Pseudo R-squ. (CS):	0.1348				
Covariance Type:	nonrobust						
	coef	std err	z	P> z	[0.025	0.975]	
Intercept	-4.2857	0.592	-7.245	0.000	-5.445	-3.126	
streams_log	0.3003	0.042	7.098	0.000	0.217	0.383	
Train Set Metrics:							
MSE :	94.9752744414452						
R2 :	0.6869834732907787						
Test Set Metrics:							
MSE :	80.58732496803977						
R2 :	0.7403063473673328						

Regression Using ElasticNet

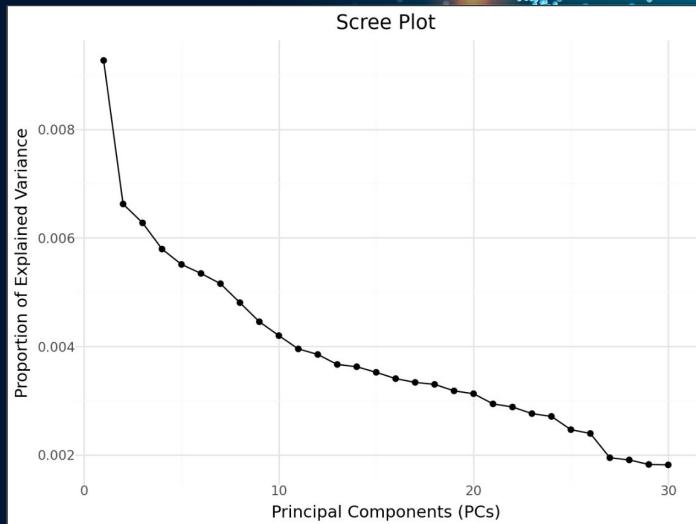
- Goal:
 - Wanted to figure out what makes a song popular and predict its popularity based on its features
- Why Linear Regression using **ElasticNet**?
 - A combination of Lasso and Ridge regression. It's great because it not only selects the most important features but also balances them, ensures no overfitting
- Shows predictive power, but it also suggests that other factors, like marketing or cultural trends, likely play a role in popularity

```
Best Parameters: {'alpha': 1.0, 'l1_ratio': 0.8}
R2 Score: 0.20117376644274243
Mean Squared Error: 225.26662559553017
```

PCA

Why PCA?

- PCA ensures the models focus on the most meaningful aspects of the data, reducing complexity
- Linear Regression predicts continuous probability, and Logistic regression predicts binary probability





Comparison of Models

Beta Regression, Clustering, ElasticNet, PCA

Beta Regression

- Test MSE of 80.58 and test R2 of 0.74 indicate passable performance with moderate error
- log features decreases skewness and can increase accuracy but also can negatively impact interpretability

Clustering

- Visually, GMM did not create good clusters
- Low Log Likelihood
 - Tried DBSCAN and KMeans but data did not do well
- Variables lowly correlated indicating clustering is not ideal



ElasticNet

R2 of .20 and MSE of 225

Leftover 80% of the variability unexplained

Limited by the quality and scope of the input features

Best for feature interpretability only

Limited predictive power

PCA

Popularity - R2 of .22 and MSE of 220

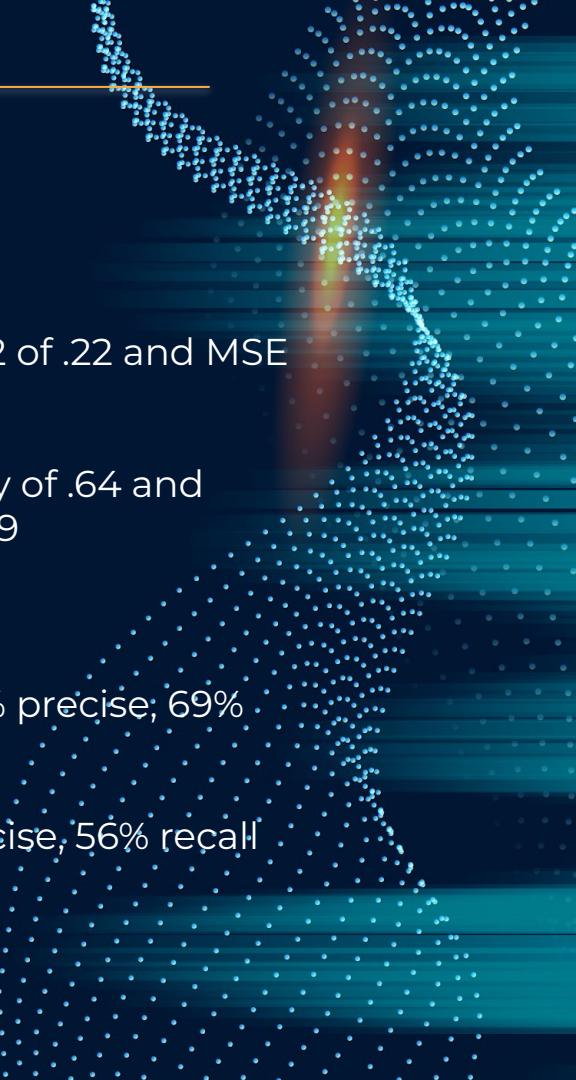
Hits - Accuracy of .64 and ROC AUC of .69

Classification:

Non-hits - 70% precise, 69% recall

Hits - 55% precise, 56% recall

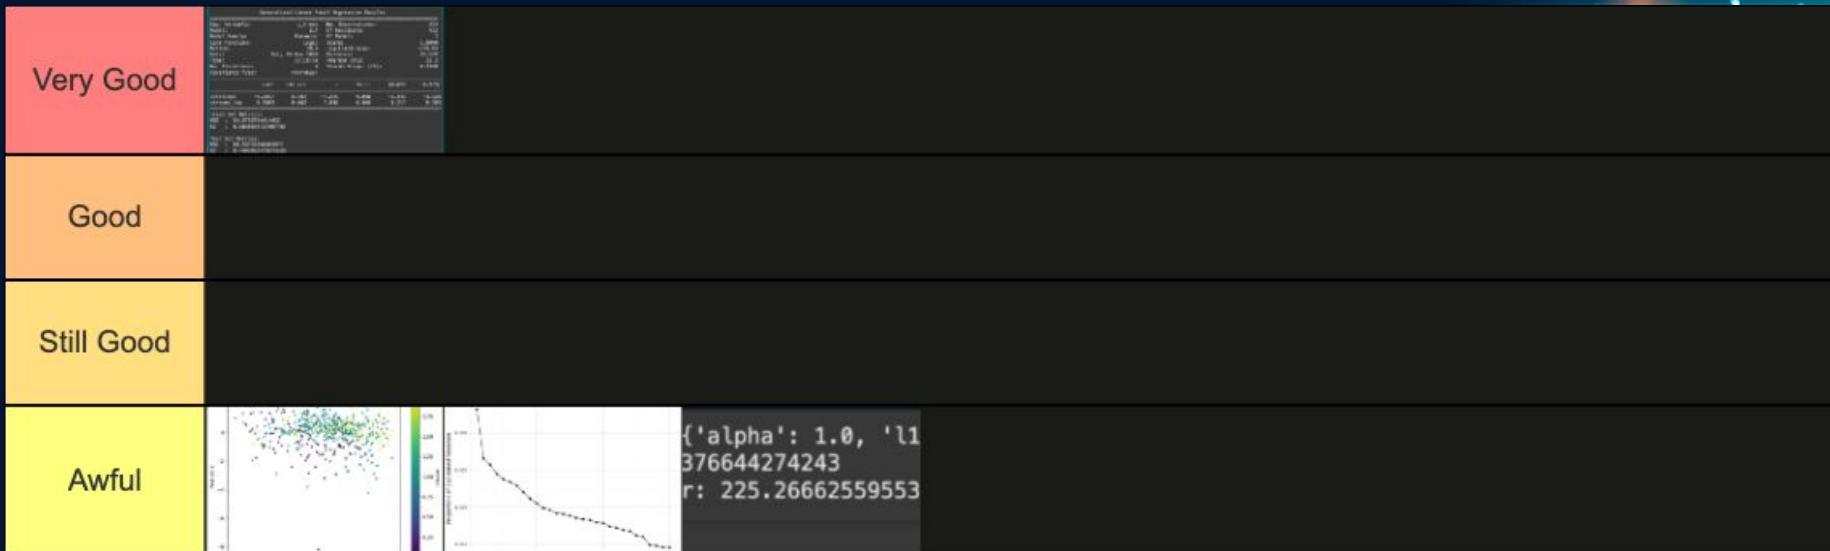
Used 30 PCs





Conclusion

Should This Model be Implemented?



Outside of Beta Regression, no model performed exceedingly well. This is displayed in the tier list above. Streams is the most important variable when predicting popularity. Features are lowly correlated.